

Part A:

Model Architecture:

```
AlexNet(  
  (features): Sequential(  
    (0): Conv2d(3, 96, kernel_size=(11, 11), stride=(4, 4))  
    (1): ReLU(inplace=True)  
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(96, 256, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (4): ReLU(inplace=True)  
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(256, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU(inplace=True)  
    (10): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU(inplace=True)  
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Dropout(p=0.5, inplace=False)  
    (1): Linear(in_features=9216, out_features=4096, bias=True)  
    (2): ReLU(inplace=True)  
    (3): Dropout(p=0.5, inplace=False)  
    (4): Linear(in_features=4096, out_features=4096, bias=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=4096, out_features=4, bias=True)  
  )  
)
```

Result for AlexNet:

[Epoch 100] train accuracy: 0.9904, loss: 0.0331

[Epoch 100] eval accuracy: 0.9015, loss: 0.5498

Part B:

Model Architecture:

```
AlexNetLargeKernel(  
  (features): Sequential(  
    (0): Conv2d(3, 96, kernel_size=(21, 21), stride=(8, 8), padding=(1, 1))  
    (1): ReLU(inplace=True)  
    (2): Conv2d(96, 256, kernel_size=(7, 7), stride=(2, 2), padding=(2, 2))  
    (3): ReLU(inplace=True)  
    (4): Conv2d(256, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (5): ReLU(inplace=True)  
    (6): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(2, 2))  
    (9): ReLU(inplace=True)  
  )  
  (classifier): Sequential(  
    (0): Dropout(p=0.5, inplace=False)  
    (1): Linear(in_features=9216, out_features=4096, bias=True)  
    (2): ReLU(inplace=True)  
    (3): Dropout(p=0.5, inplace=False)  
    (4): Linear(in_features=4096, out_features=4096, bias=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=4096, out_features=4, bias=True)  
  )  
)
```

Result for AlexNetLargeKernel:

[Epoch 100] train accuracy: 0.9937, loss: 0.0175

[Epoch 100] eval accuracy: 0.8724, loss: 1.1351

Model Architecture:

```
AlexNetAvgPooling(  
  (features): Sequential(  
    (0): Conv2d(3, 96, kernel_size=(11, 11), stride=(4, 4))  
    (1): ReLU(inplace=True)  
    (2): AvgPool2d(kernel_size=3, stride=2, padding=0)  
    (3): Conv2d(96, 256, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (4): ReLU(inplace=True)  
    (5): AvgPool2d(kernel_size=3, stride=2, padding=0)  
    (6): Conv2d(256, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU(inplace=True)  
    (10): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU(inplace=True)  
    (12): AvgPool2d(kernel_size=3, stride=2, padding=0)  
  )  
  (classifier): Sequential(  
    (0): Dropout(p=0.5, inplace=False)  
    (1): Linear(in_features=9216, out_features=4096, bias=True)  
    (2): ReLU(inplace=True)  
    (3): Dropout(p=0.5, inplace=False)  
    (4): Linear(in_features=4096, out_features=4096, bias=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=4096, out_features=4, bias=True)  
  )  
)
```


Result for AlexNetAvgPooling:

[Epoch 100] train accuracy: 0.9958, loss: 0.0141


[Epoch 100] eval accuracy: 0.8838, loss: 0.7957

Part C: Domain Kernels:

Conv2d-0



Conv2d-3



Conv2d-6




Conv2d-8



First 12 Out-Channels
First 12 In-Channels

Conv2d-10




Class Kernels:

Conv2d-0



Conv2d-3



Conv2d-6




Conv2d-8



First 12 Out-Channels
First 12 In-Channels

Conv2d-10



My findings:

The kernels start as simple and then progress to more advanced convolutions. There are also some consistencies in both domain and category classification even though they are looking for different properties in the images.

